

RJava User and Developer Manual

Yi Lin
yi.lin@anu.edu.au

November 4, 2013

Abstract

RJava is a restricted subset of the Java language with low-level extensions that allow access to hardware and operating system. RJava utilizes the same syntax as Java, and consequently inherits benefits from Java such as type safety, various software engineering tools and productivity. Furthermore, by restrictions, RJava is a fully static language with closed world assumption. Thus it requires a much more succinct runtime, and is well suitable for aggressive static compilation and optimizations. RJava is designed to be an implementation language for virtual machine construction (and more broadly for system programming).

This manual describes the language and its current implementation—the RJava Compiler (RJC). It is intended for RJava users and developers who are willing to contribute. This manual will be maintained to keep pace with the RJC code base.

Contents

1	RJava Basics	3
1.1	RJava Restrictions	3
1.1.1	Restriction Rules and Rulesets	3
1.1.2	Restriction Model	4
1.1.3	RJavaCore Ruleset	5
1.1.4	Other Predefined Restrictions	6
1.2	RJava Extensions	7
1.2.1	org.vmmagic	7
1.2.2	org.rjava.osext	8
1.3	Relation between RJava and MMTk/JikesRVM	9
2	RJava Compiler Tools	11
2.1	Command Line Options	11
3	RJava Compiler Implementation	11
3.1	Codebase Overview	11
3.2	Basic Workflow	11
3.3	Unit Tests	11
4	RJava Compiler Details	11
4.1	Magic/Unboxed Types	11
4.2	java.lang.* Package	11
4.3	RJava Compiler AST	11
4.4	Analysis and Optimization passes	11
4.5	RJava Compiler Targets	11
5	MMTk/RJava Manual	11
5.1	Unofficial Changes	11
5.2	MMTk-VM Interface	11

1 RJava Basics

This section describes about RJava's restrictions and extensions, which make RJava differentiate from the vanilla Java.

What is RJava?

Generally speaking, unless what is *restricted* and what is *extended*, the other part of RJava remains the same as Java (syntactically and semantically) . Besides, the restrictions and the extensions are also compliant with Java *syntax*.

Thus, Java compilers and static analyzers can parse RJava code (they are syntactically same). Java runtimes can execute RJava programs as well; However, Java runtimes will not be able to recognize RJava specific restrictions and extension, thus executing RJava with Java runtimes may result in wrong results, which is not encouraged.

RJava is also compliant with existing Java editors and IDEs. When using Java IDE to edit RJava code, extensions and restrictions need to be imported and also try avoid using execution from IDE.

1.1 RJava Restrictions

One important feature about RJava is that all its restrictions are formalized and any RJava code need to declare the restrictions it complies with. The benefit is that both the developers and the RJava compilers will be able to see the declaration and ensure the restrictions are met.

1.1.1 Restriction Rules and Rulesets

RJava restrictions are defined as Java annotations, and they are annotated by `@RestrictionRule` (`org.rjava.restriction.RestrictionRule`). The following is how `@NoException` restriction is defined:

```
@RestrictionRule
public @interface NoException{
}
```

Restriction rules can be used to annotate classes or methods, in the same way as Java annotations are used. However, it is *not* encouraged to directly

use restriction rules. The preferred way is to form a `@RestrictionRuleset` by a set of chosen `@RestrictionRule`, and annotate a certain scope of code with such `@RestrictionRuleset`.

One example is the MMTk scope and its restriction `@MMTk`. The following ruleset indicates that any code within `@MMTk` needs to obey `@RJavaCore` ruleset (any RJava code need to include this ruleset, it will be described in next subsection) as well as two additional restrictions (`@NoRuntimeAllocation` and `@Uninterruptible`). The declaration is as below:

```
@RestrictionRuleset

@RJavaCore
@NoRuntimeAllocation
@Uninterruptible
public @interface MMTk{
}
```

1.1.2 Restriction Model

Restrictions follow those rules:

- Restriction rules and restriction rulesets apply to scopes of classes and methods to declare restrictions, *not* to fields.
- Restriction rules and rulesets may apply to another restriction ruleset, to indicate that the latter ruleset *includes* those rules and rulesets. This may happen iteratively.
- Any RJava class needs to be restricted by `@RJavaCore`, or restricted by a ruleset that (iteratively) includes `@RJavaCore`.
- A scope restricted by a restriction ruleset is restricted by every single restriction rule within that ruleset.
- If a class is restricted by Rule A, all its methods are restricted by Rule A unless the method is restricted otherwise (particularly `@Uninterruptible` vs. `@Interruptible`).
- Restrictions on a class do *not* affect its child classes. Restrictions on a method do *not* affect its overriding methods.

1.1.3 RJavaCore Ruleset

@RJavaCore declares the basic restrictions that any RJava-compliant code need to follow. It declares as:

```
@RestrictionRuleset

@NoDynamicLoading
@NoReflection
@NoException
@NoCastOnMagicType
@NoExplicitLibrary
@NoEnum
public @interface RJavaCore {
}
```

Specifically,

- @NoDynamicLoading: the code is not allowed to dynamically load classes. To forbid this, `java.lang.Class.forName()` and `java.lang.Classloader.loadClass()` are not allowed. Also inheriting from `java.lang.Classloader` is also not allowed.
- @NoReflection: the code is not allowed to use any reflection-based feature. To forbid this, any method that may return `java.lang.Class` is not allowed. @NoExplicitLibrary already forbids the use of `java.lang.reflect` package. Furthermore, `java.lang.Object.getClass()` is not allowed (`.class` syntax is also not allowed, since on bytecode level, it gets translated into `Object.getClass()`.)
- @NoException: the code is not allowed to throw exceptions. To forbid this, `throws` in method declarations and also `throw` statement are not allowed. Catching blocks are allowed in order to maintain correct syntax when using library methods, however, there will be *no* exception thrown and catching blocks will *not* get executed in any case. It is preferred to write such code as below:

```
Object lock = new Object();
try {
    lock.wait();
} catch (InterruptedException ignore) {}
```

- **@NoCastOnMagicType**: the code is not allowed to do any type casting when *at least one* side is RJava magic types. RJava magic types reside in `org.vmmagic.unboxed` package, including 5 unboxed magic types (`Address`, `Extent`, `ObjectReference`, `Offset`, `Word`) and their array counterparts. Allowed type casting for magic types can be done via provided methods in those classes (more will be discussed in the section about extensions).
- **@NoExplicitLibrary**: the code is not allowed to import any java library. However, classes in the `java.lang.*` package are implicitly imported to any Java code, and they are entangled with Java syntax. RJava *allows* the use of `java.lang.*`. Current RJava Compiler has only implemented a subset of the `java.lang.*` package, more complete implementation will be done in the future.
- **@NoEnum**: the code is not allowed to use enumerate type. Enumerate type is forbidden since enabling full features of Java `enum` would highly involve with dynamic behaviors and extensive use of library methods from `java.lang.Enum`. We designed RJava to be a *fully static* language, thus we forbid the use of enumerate type.

These core restrictions define the RJava language. Any valid RJava code needs to declare their compliance with **@RJavaCore**.

By these restrictions, RJava 1) allows closed world assumption, which helps aggressive static compilation, 2) supports little dynamic behaviors and bares very succinct execution runtime, which makes RJava possible to run with limited hardware resources, and 3) is a simple yet still expressive language.

1.1.4 Other Predefined Restrictions

RJava defines the above *restriction model*. RJava encourages its users to utilize this model and, if favorable, add their own restriction rules and rule-sets to more precisely describe restrictions on their own scope. However, RJava predefined a set of restrictions. Some are already mentioned in the sections above. These predefined restrictions are subject to change, check `org.rjava.restriction.rules` package in the code base for the latest information.

1.2 RJava Extensions

In order to undertake system programming task, RJava introduces extensions to allow efficient access to hardware and operating systems. RJava extensions include 1) the `org.vmmagic` package (by Daniel et al.) that allows memory/address representation and operations, and 2) the `org.rjava.osext` package that provides access to operating systems (including some system calls).

1.2.1 `org.vmmagic`

The `org.vmmagic` package provides 5 unboxed magic types to describe pointer-like types and their operations. Though implementation may vary, ‘unboxed’ means those types are not normal RJava objects, and they are more like primitive types though there are methods declared for each type. These magic types share the same length as the pointer length on the target machine:

- **ObjectReference**: can be cast from and to an RJava object. *No* arithmetic, load/store or comparison operations are supported. **ObjectReference** can be cast to **Address** for further operations (unsafe).
- **Address**: used as a pointer-alike type. Arithmetic, load/store and comparison operations are supported.
- **Extent**: used to describe size in bytes (unsigned positive value) Arithmetic and comparison operations are supported.
- **Offset**: used to describe offset in bytes (signed). Arithmetic and comparison operations are supported.
- **Word**: used as a pointer-sized integer. Arithmetic, comparison and bit-wised operations are supported. Can be cast from and to the other 4 magic types.

The above unboxed magic types each have an array counterpart, namely **ObjectReferenceArray**, **AddressArray**, **ExtentArray**, **OffsetArray**, **WordArray**. They provide `create()`, `get()/set()` and `length()` operations. Arrays of those magic types should *only* be created by using these types.

Besides magic types, `org.vmmagic` provides ‘pragma’ for supplying information to the compiler. Useful ‘pragmas’ include `@Inline`, `@NoInline`, `@NoBoundsCheck`, `@NonNullCheck`, etc. Some ‘pragmas’ from `org.vmmagic` are now considered as an RJava `@RestrictionRule` such as `@Uninterruptible`, and some are very specific to Java and not applicable for RJava thus will be deleted. This part is still a draft, and subject to change.

1.2.2 `org.rjava.osext`

The `osext` extension allows access to operating systems. Currently `osext` only includes a minimum set of methods for implementing a memory manager (MMTk), and it *will* expand during further development.

- `OSConcurrency`: concurrency/threading related methods.
 - `void mutexLock/Unlock(Object lock)`: provides an alternative to using synchronization on an RJava object.
 - `void threadSuspend/Resume(Thread t)`: suspending an RJava thread (unsafe)
- `OSMemory`: memory related methods.
 - `Address malloc(int size)`: allocates raw memory
 - `Address mmap(Address start, Extent length, int protection, int flags, int fd, Offset offset)`: calls `mmap` system call
 - `int mprotect(Address start, Extent length, int prot)`: calls `mprotect` system call
 - `Address memset(Address start, int c, Extent length)`: calls `memset` system call
- `OSNative`: other methods that calls to native C
 - `int errno()`: a wrap of `errno()` in `errno.h`
 - `String strErrno()`: a wrap of `strerror()`
 - `double random()`: generates a random double between 0 and 1

1.3 Relation between RJava and MMTk/JikesRVM

The RJava project was motivated by trying improve the portability of Memory Management ToolKit (MMTk). MMTk started with JikesRVM project, and serves as its memory manager. The same as the rest part of JikesRVM, MMTk is also written in a variant of Java (ad-hoc restrictions with `org.vmmagic` extensions). Though MMTk was designed to be a portable language-agnostic memory manager, its portability was never a success.

One of the main reasons that constrain its portability is the portability of the language it is written in. On one hand, such Java variant is specially tailored, and requires support from its hosting runtime. It is no longer a ‘write once run everywhere’ Java program, and it cannot execute on a stock Java VM. Whoever wants to host MMTk needs to find a way to execute the Java variant that MMTk is written in. On the other hand, MMTk is written in such an Java variant, and there is an inneglectable performance cost to integrate MMTk(Java) with a hosting runtime written in C/C++ which is the most common case. Past experiences (VMKit, Rotor, GHC) of porting MMTk took an approach of ahead-of-time (AOT) compiling MMTk to native codes, which overcame these two obstacles. However, *none* of their approaches is general or reusable, and each of them took a great effort. Furthermore, in term of performance, *none* of the experiences was reported as a success (meanwhile MMTk hosted by JikesRVM achieves excellent performance).

Thus the idea of formalizing the Java variant that MMTk is written in and providing an AOT compiler to lower such language for effective integration with low-level C/C++ code becomes a promising solution to MMTk portability issues. This is one part of the motivations.

Another important motivation that evolves the RJava project is the clear benefits from using a higher-level language for virtual machine implementation. Compared with traditional approach of using C/C++, higher-level languages provide benefits in safety and productivity both of which are valuable in VM construction. However, existing experiences of using high-level languages for VM construction (including JikesRVM and MMTk) all formed their own variant of such high-level languages, which not only results in inability of code reusing but also practically introduces metacircularity issues. As we already found out, metacircularity should be orthogonal to using high-level languages; however, in practice so far, whenever high-level languages are used in virtual machine construction, the chosen implementation language is

always the target language of the VM. Thus this project not only wants to settle the reusability/portability issues but also wants to create a general implementation language that can deliver high-level language benefits without falling into metacircular traps.

In short, RJava originates from the ad-hoc coding pattern used in MMTk, and absorbs the `org.vmmagic` package which is also used in MMTk. However, it evolves to be a more general implementation languages to deliver higher-level language benefits as well as good performance.

- 2 RJava Compiler Tools**
 - 2.1 Command Line Options
- 3 RJava Compiler Implementation**
 - 3.1 Codebase Overview
 - 3.2 Basic Workflow
 - 3.3 Unit Tests
- 4 RJava Compiler Details**
 - 4.1 Magic/Unboxed Types
 - 4.2 java.lang.* Package
 - 4.3 RJava Compiler AST
 - 4.4 Analysis and Optimization passes
 - 4.5 RJava Compiler Targets
- 5 MMTk/RJava Manual**
 - 5.1 Unofficial Changes
 - 5.2 MMTk-VM Interface